

Big Data - Labo

Tuur Vanhoutte

22 februari 2021

Inhoudsopgave

1	Intro	1
2	NAT-ing	1
2.1	NAT	1
2.1.1	The problem	1
2.1.2	The solution	2
2.2	SSH Tunnel	2
3	Container technology	2
3.1	Docker	2
3.2	Microservices	3
3.2.1	Monolithic vs Microservices	3
3.3	Virtualization vs Containerization	4
3.3.1	Virtualization	4
3.3.2	Containerization	4
3.4	Shared kernel	4
3.4.1	What is a kernel?	4
3.4.2	How?	6
3.4.3	Namespaces	6
3.4.4	Containers	6
3.5	Images	6
3.5.1	Image layer	7
3.6	Docker is lightweight	7
3.7	Using Docker	8
3.7.1	Layers bekijken	8
3.7.2	Make Dockerfile more efficient	8
3.7.3	Connecting to a database in a different container	8

1 Intro

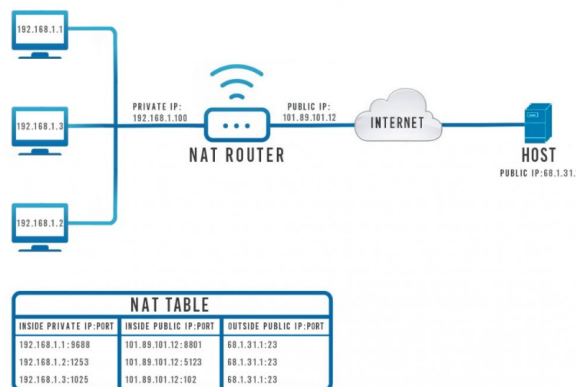
Topics:

- Linux basics + containers
- Elastic search (text search, document store)
- Linux Batch Processing & Dask
- InfluxDB (timeseries)
- Cloud services (Kafka, Kinesis, Lambda, ML services, ...)

2 NAT-ing

2.1 NAT

= Network Address Translation



Figuur 1: NAT diagram

2.1.1 The problem

- We only have one (public/private) IP-address
 - Howest: 172.23.82.60
- Connecting to a server over a network:
 - Using a protocol (HTTP) which uses TCP
 - Our server has an IP address: 172.23.82.60
 - Our server is listening at port 5000
 - \Rightarrow `http://172.23.82.60:5000`
- Problem: We want to have multiple IP addresses
 - Student 1 wants to reach `http://192.168.20.21:5000`
 - Student 2 wants to reach `http://192.168.20.22:5000`
 - Student x wants to reach `http://192.168.20.xx:5000`

2.1.2 The solution

Translation is needed!

- 172.23.82.60:5000 should point to 192.168.20.21:5000
- 172.23.82.60:5001 should point to 192.168.20.22:5000
- 172.23.82.60:5xxx should point to 192.168.20.xx:5000

We can use any port, on both sides:

- 172.23.82.60:8000 can point to 192.168.20.21:5000
- 172.23.82.60:8000 can point to 192.168.20.21:3000

2.2 SSH Tunnel

= SSH Port Forwarding

Resource	Internal IP	Username	Password	External port	Internal port
Vyos Router	192.168.50.1	vynos	P@ssw0rd	7000	22
Storage	192.168.50.2	student	P@ssword	n.v.t.	22
SSH	192.168.50.3	student	P@ssword	7040	22
RDP	192.168.50.4	Administrator	P@ssword	7020	3389
vCenter vSphere	192.168.50.10	administrator@vsphere.local	P@ssword	7060	443
vCenter appliance	192.168.50.10	root	P@ssword	n.v.t.	5480
ESXi-00	192.168.50.11	root	P@ssword	n.v.t.	22
ESXi-01	192.168.50.12	root	P@ssword	n.v.t.	22

Figuur 2: Example



Figuur 3: Voorbeeld: een tunnel wordt opengemaakt en er wordt ingelogd in user@instance

3 Container technology

3.1 Docker

- Docker = ecosystem for creating and running containers
- Docker wants to make it possible to install and run software on any system

- Other reasons: Microservices/DevOps/Resource usage
- Docker != Container
 - Docker CLI
 - Docker Engine
 - Docker Image
 - Docker Container
 - Docker Hub
 - Docker Compose
 - Docker Swarm
 - ...

3.2 Microservices

- = A software development technique
- Structure an application as a collection of loosely coupled services
- Lightweight
- Microservices-based architectures enable continuous delivery and deployment
- <https://en.wikipedia.org/wiki/Microservices>

3.2.1 Monolithic vs Microservices

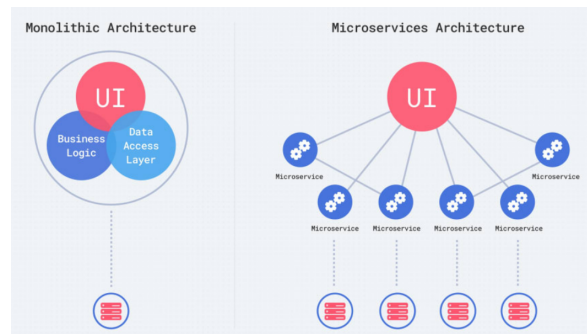


Figure 4: Monolithic architecture vs Microservices architecture

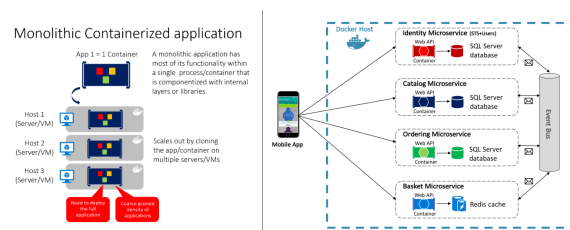
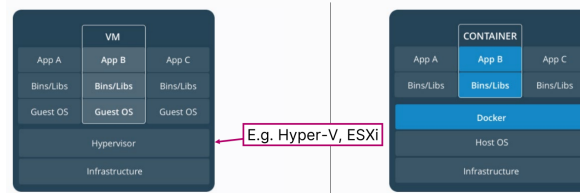


Figure 5: Monolithic Containerized application

Microservices does **not** necessarily mean containerization!

3.3 Virtualization vs Containerization



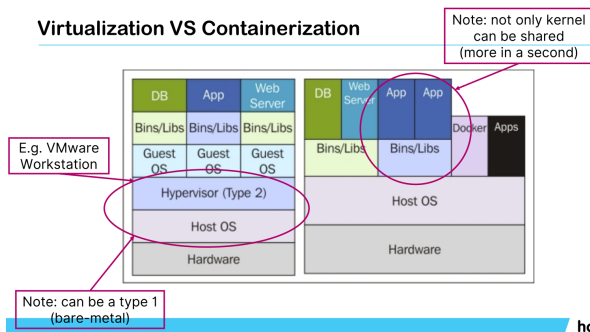
Figuur 6: Virtualization vs Containerization

3.3.1 Virtualization

- = An abstraction of physical hardware turning one server into many servers
- Multiple VMs can run on the same machine
- Each VM includes a full copy of an Operating System (OS), one or more apps
- Takes a lot of space
- Can be slow to boot

3.3.2 Containerization

- = An abstraction at the app layer that packages code and dependencies together
- Multiple containers can run on the same machine, they share the OS kernel with each other, each running as isolated processes in user space.
- Takes up less space than VMs
- Boot up almost instantly



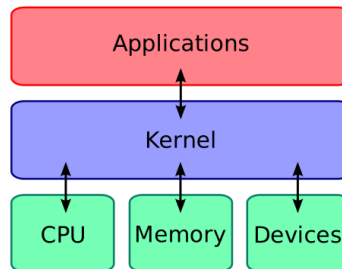
Figuur 7: Schematic

3.4 Shared kernel

3.4.1 What is a kernel?

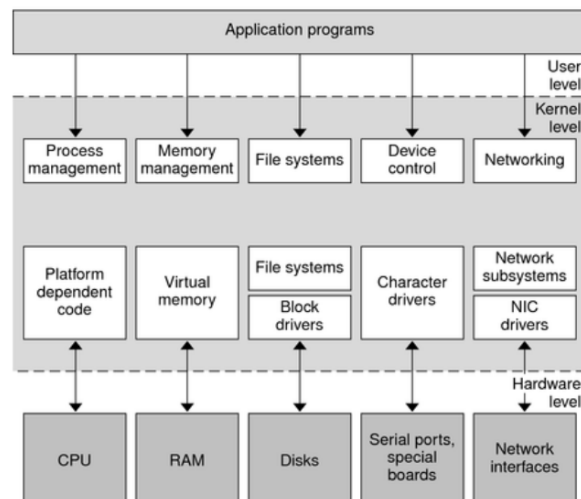
- Piece of software that offers basic functionality to the OS

- System calls: open, read, write, close, wait, exit, . . .
- A typical kernel has a few hundred system calls



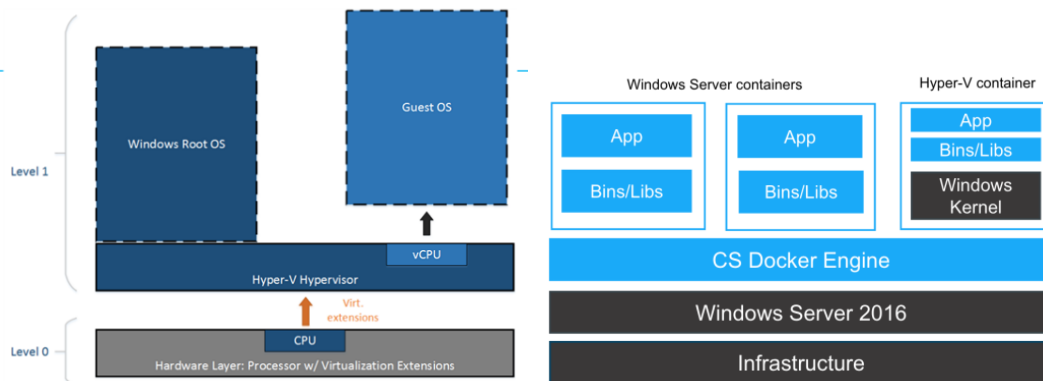
Figuur 8: The kernel is the layer that communicates between hardware and applications

- Docker shares the host OS kernel
 - Host OS: Windows / MacOS / Linux
 - Shared Linux Kernel



Figuur 9: Kernel in detail

- The Ubuntu container requires the Linux kernel
- The Linux kernel runs in a Virtual Machine



Figuur 10

3.4.2 How?

Two important Linux kernel features:

- **Namespaces** are a feature of the Linux kernel that partitions kernel resources
- **cgroups** (control groups) is a Linux kernel feature that limits, accounts for, and isolates resource usage of a collection of processes

Simpler:

- Namespaces = isolating resources per process (or group of processes)
- cgroups = Limiting resource usage per process (or group of processes)

3.4.3 Namespaces

- 7 types:
 - mount, UTS, IPC, network, PID, cgroup, user
- For the process (or group of processes) it looks like there is a completely isolated set of resources

3.4.4 Containers

What is a container?

- One or more running processes (if not running anymore \Rightarrow container dead)
- Resources are specifically assigned to it
- The real building blocks: Linux kernel features
 - Namespaces
 - cgroups

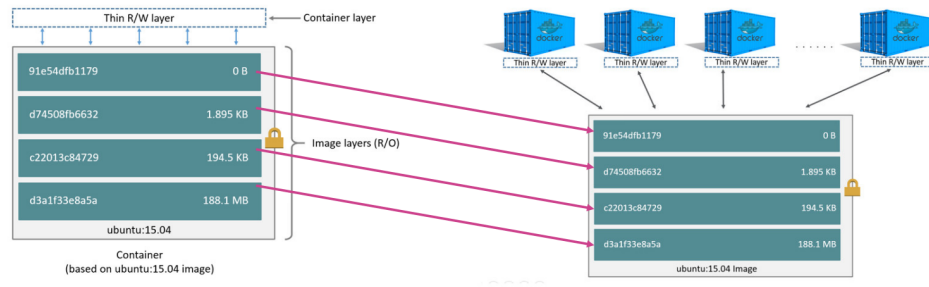
3.5 Images

What is an image?

- Filesystem snapshot
- Startup command
- Layered structure (!)

Instance of image = container

3.5.1 Image layer



Figuur 11: Image layers

- RUN, COPY, ADD
 - = new read-only layer
- Top layer = container layer
 - Writeable
- Delete container = delete container layer
 - Image will still exist
 - Persistent volumes

3.6 Docker is lightweight

- Shared kernel
- Container has no OS
- Less disk space ⇒ sharing layers
- Small community images
 - ex: Alpine Linux (small, simple, secure)
- Current Docker version is using runC (previously LXC = Linux Containers)
 - runC = tooling (written in Go) that makes it possible to create and run containers
 - runC = CLI to 'easily' access kernel features such as cgroups and namespacing
 - runC = successor of libcontainer (developed by Docker)
 - Open-sourced ⇒ better community
 - runC implements 'Open Container Initiative Runtime Specification'

– <https://github.com/opencontainers/runtime-spec>

Docker is ‘nothing more’ than an ecosystem about creating & running containers

3.7 Using Docker

(see slides 40-55 in [02_big_data_01_containers.pdf](#) for basic commands)

3.7.1 Layers bekijken

With the command ‘docker history <image | container id>’ you’ll get an overview of the layers of an image.

- Every RUN, COPY, ADD adds a new read-only layer
- Make Dockerfile more efficient ⇒ create less layers

3.7.2 Make Dockerfile more efficient

Our Dockerfile, before optimisation:

```
1 FROM python:3.9.1-alpine3.13
2 WORKDIR '/app'
3 RUN apk add --no-cache linux-headers g++
4 RUN pip install Flask # we can replace these two lines by:
5 RUN pip install uwsgi # RUN pip install -r requirements.txt
6 COPY ./ ./
7 RUN addgroup -S uwsgi && adduser -S uwsgi -G uwsgi
8 USER uwsgi
9 CMD ["uwsgi", "--ini", "app.ini"]
```

After optimisation:

```
1 FROM python:3.9.1-alpine3.13
2 WORKDIR '/app'
3 RUN apk add --no-cache linux-headers g++
4 # the addgroup and adduser commands can be higher up
5 RUN addgroup -S uwsgi && adduser -S uwsgi -G uwsgi
6 # first, we copy the requirements.txt file
7 COPY ./requirements.txt ./
8 # then we install ALL packages
9 RUN pip install -r requirements.txt
10 # then we copy the remaining files
11 COPY ./ ./
12 USER uwsgi
13 CMD ["uwsgi", "--ini", "app.ini"]
```

3.7.3 Connecting to a database in a different container

Use ‘ip a’ to find the correct ip to use in this command:

```
1 docker run -p 8080:8080
2 -e POSTGRES_PASSWORD=student_password
```

```
3 -e POSTGRES_USER=student_user
4 -e POSTGRES_DATABASE=labo
5 -e POSTGRES_PORT=5432
6 -e POSTGRES_HOST=ip-van-je-vm # change this ip
7 -e PORT=8080
8 jouw-naam/api # change this
```