

Linux OS

Tuur Vanhoutte

March 6, 2021

Contents

1	Introductie	1
1.1	Verschil Server & Workstation	1
1.1.1	Server	1
1.1.2	Workstation	1
1.2	Extra information/resources	1
1.3	What is Linux?	1
1.3.1	What is an operating system (OS)?	1
1.3.2	What is a Kernel?	2
1.4	GNU Operating System	2
1.5	Linux, the kernel	2
1.5.1	Distributions	2
1.6	Open Source	3
1.6.1	Commercial distributions	3
1.6.2	In this course: Debian	3
2	Debian Installation	4
2.1	Networking in Linux (with VMWare)	4
2.2	Users in Linux	4
2.3	Disks, partition, filesystems	4
2.3.1	Partitions	5
2.4	MBR <> GPT	5
2.4.1	MBR	5
2.4.2	GPT	6
2.4.3	Bootstrap procedure	6
2.4.4	Linux boot process	7
2.4.5	BIOS <> UEFI	7
2.5	Filesystems	7
2.5.1	Windows	7
2.5.2	Linux	7
2.5.3	Swap	8
2.6	File structure	8
2.7	Configuration	9
2.7.1	Packages	9
2.7.2	Package management	9
2.7.3	Useful packages	10
2.8	Shutdown of VM	10
2.9	Basic network	10
2.9.1	Basic networking commands	11
2.10	Services	11
2.11	Wooclap Questions	11
3	File structure	12
3.1	Intermezzo: single user mode	12
3.1.1	Runlevels	13
3.2	Intermezzo: Add disk	13
3.2.1	What after a reboot?	14
3.3	Navigate through the tree	14
3.3.1	Relative vs absolute path	14
3.4	Filesystem Hierarchy Standard (FHS)	14
3.4.1	Rules in the standard	15

3.5	Some useful tips	17
3.5.1	History	17
3.5.2	Bind mount	17
3.5.3	dd	17
3.6	Wooclap Questions	17
4	Filesystems	18
4.1	Introduction	18
4.2	Blocks	18
4.3	ext2/3/4	19
4.3.1	Journaling	19
4.4	RAID	19
4.4.1	RAID Controller	19
4.4.2	RAID 0	19
4.4.3	RAID 1	20
4.4.4	RAID 4	20
4.4.5	RAID 5	21
4.4.6	RAID 6	21
4.4.7	Disk failure	21
4.4.8	Compound RAID levels	22
4.5	OpenZFS	22
4.5.1	ZFS	22
4.5.2	Open ZFS	22
4.6	Intermezzo: Kernel modules	22
4.6.1	Commands	23
4.7	Intermezzo: Snapshots	23
4.7.1	Do we still need backups if we have snapshots?	23
5	File manipulation	23
5.1	Basics	23
5.2	Bundle files	24
5.3	Links and inodes	24
5.3.1	Inodes	24
5.3.2	Symbolic links	25
5.3.3	Hardlinks	25
5.4	File permissions	25
5.5	Overview of basic commands	26
5.6	Wooclap	27
6	Text editors, Piping, Redirection & Jobs	27
6.1	Text editors	27
6.1.1	vi vs vi-improved	28
6.1.2	First steps in vim	28
6.1.3	Search and replace	29
6.1.4	Basic editing tricks	29
6.2	Piping	30
6.3	Redirection	30
6.3.1	stdout and stderr	30
6.3.2	stdin	31
6.4	Jobs and process Management	31
6.4.1	Exit codes	32
6.4.2	Combining commands	32

6.4.3	Jobs	33
6.4.4	Inter-process Communication	33
6.5	Intermezzo: System Load	33
6.6	Some useful tips	34
6.6.1	With which unique IP-addresses are there open sockets and how many?	34
6.6.2	TTY	34

1 Introductie

1.1 Verschil Server & Workstation

1.1.1 Server

- Deliver services to (multiple) users
- Focussed: only this and nothing else
- Secure
- No GUI, everything happens through the commandline
- ⇒ as small a footprint as possible

1.1.2 Workstation

- Use services
- Create documents
- Look for information
- Consume multimedia
- GUI
- ⇒ Large footprint

1.2 Extra information/resources

- The Linux Documentation Project: <http://tldp.org>
- Pluralsight LPIC-1: Linux Professional Institute Certification: <https://www.pluralsight.com/paths/lpic-1>
- The Arch Linux Wiki is one of the most extensive sources of info about Linux: <https://wiki.archlinux.org>
 - In this module we will use Debian, not Arch, but many things are very similar
- Google

1.3 What is Linux?

1.3.1 What is an operating system (OS)?

Definitie 1.1 (Operating System) *An operating system, or OS, is software that communicates with the hardware and allows other programs to run.*

It is comprised of system software = the fundamental files your computer needs to function.

Linux is NOT an operating system: Linux = the kernel

1.3.2 What is a Kernel?

Definitie 1.2 (Kernel) *The kernel is software that is the core of a computer's operating system, with complete control over the system.*

It is the first program loaded on start-up.

It handles...:

- ... the rest of the startup
- ... input/output requests from software, translating them into instructions for the CPU
- ... memory
- ... peripherals

1.4 GNU Operating System

Definitie 1.3 (GNU) *GNU = GNU's Not Unix (recursive algorithm)*

Founded by Richard Stallman (ex-MIT, founder of the Free Software Foundation), 1984

Goal: completely free Operating System

1.5 Linux, the kernel

By Linus Torvalds (Finland), 1991

- Own personal development, not initially intended to distribute
- Interest from other developers, mainly to use with GNU OS
- Meanwhile contributions of over 12000+ developers
- 492 of top-500 supercomputers in the world run Linux
- Basis for Android, Chrome OS

Linux = the kernel

GNU = OS-tools around the kernel

⇒ **GNU/Linux**


1.5.1 Distributions

Definitie 1.4 (Distribution) *A Linux distribution (or distro for short) is GNU/Linux + extra tools and applications to create a full-fledged OS.*

That distribution can be easily copied and installed to other computers.

- RedHat (CentOS)
- Debian (Ubuntu)
- Arch Linux
- Void Linux
- Gentoo
- Pop! OS

Red Hat family tree



Debian family tree




Figure 1: https://upload.wikimedia.org/wikipedia/commons/1/1b/Linux_Distribution_Timeline.svg

https://en.wikipedia.org/wiki/List_of_Linux_distributions

1.6 Open Source

Definitie 1.5 (Open Source) *Open source software is software of which the code is licensed to be open to everyone.*

Anyone can use, change, distribute the software. This allows code to be developed in a public manner.

OPEN SOURCE DOES NOT MEAN FREE

1.6.1 Commercial distributions

= Open source, non-free distributions

- SUSE Linux Enterprise Server (SLES)
- SUSE Linux Enterprise Desktop (SLED)
- Red Hat Enterprise Linux (RHEL)
- Oracle Enterprise Linux

Commercial distributions have official support channels.

⇒ You're not paying for the operating system, you're paying for the support.

1.6.2 In this course: Debian

- Current version: 10.7
- Forms the basis of many others: Ubuntu, Raspbian, Knoppix, Linux Mint
- Available on many platforms: Intel x86, AMD64, Intel64, ARM, MIPS, Power Systems, ...

2 Debian Installation

See Labs for detailed Installation tutorial

2.1 Networking in Linux (with VMWare)

- VMWare presents ethernet adapter
- During creation of virtual machine: MAC-address is created
- During installation: network configuration through DHCP
 - IPv4-address
 - Default gateway
 - DNS-server
 - Optional: proxy-server

2.2 Users in Linux

- Linux is multi-user from the ground up
 - Multiple users can be active at the same time
- 'Administrator'-user is called root
- Each user has a user-ID (uid)
 - root has uid=0
 - uid=0 has all rights
- Each user has a home-directory

2.3 Disks, partition, filesystems

- Our VM has 1 disk
 - Presented on the SCSI-bus
 - First disk on SCSI-bus: **sda**
 - Then sdb, sdc, ...
- Disk = concatenation of blocks
- Divide blocks in collections (=partitions)
 - 1st partition: sda1
 - 2nd partition: sda2
 - ...
- 2 types of partitions
 - Primary
 - Extended

2.3.1 Partitions

Primary partition

- A filesystem can be created inside this
- Up to 4 primary partitions

Extended Partition

- 'Logical' partitions can be created inside this

Our setup:

- sda1: primary partition
- sda2: extended partition
- sda5: 'logical' partition inside extended partition sda2



Figure 2: Our setup

2.4 MBR <> GPT

2.4.1 MBR

We use the MBR Partitioning scheme

Definitie 2.1 (MBR) *MBR, or Master Boot Record, is a special type of boot sector at the start of a disk.*

It contains:

- *a set of instructions necessary to boot operating systems.*
- *info about how partitions are placed on disk*

Limitations:

- Maximum disks of 2TB
- 32-bit for number of logical sectors
- Common sector size: 512 bytes
- $2^{32} \cdot 512 \text{ bytes} = 4294967296 \cdot 512 \text{ bytes} \approx 2\text{TB}$
- Maximum amount of primary partitions = 4

BIOS can boot from a disk with MBR partitioning

2.4.2 GPT

Definitie 2.2 (GPT) *GPT, or GUID Partition Table, is a standard for the layout of partition tables on a disk. It's an alternative to MBR.*

It uses unique identifiers (GUIDs)

- BIOS cannot boot from a disk with GPT-partitioning: UEFI required when using GPT
- GPT allows disks larger than 2TB

Definitie 2.3 (UEFI) *UEFI, or Unified Extensible Firmware Interface, is a newer firmware interface by Intel (90's) that replaces the BIOS interface by IBM (70's).*

How does it work?

- Disk = collection of blocks
- Group of blocks together = sector
- Common sector size: 512 bytes
- Sectors indicated with Logical Block Addresses (LBA)
- MBR in LBA 0
- GPT headers in LBA 1
- Partition tabel right after that

2.4.3 Bootstrap procedure

1. Motherboard gets electricity
2. Mini-loader hardcoded in memory
 - BIOS gets loaded
3. Boot media are consulted
4. First boot medium, first sectors are being read \Rightarrow
5. MBR contains a bit-more-advanced loader: GRUB
 - GRand Unified Bootloader
6. This loader loads a more advanced loader (GRUB second stage bootloader)
7. The OS is loaded

2.4.4 Linux boot process

6 high level steps

- BIOS (Basic Input/Output System) - loads MBR
- MBR (Master Boot Record) - loads GRUB
- GRUB (Grand Unified Bootloader) - loads kernel
- Kernel - executes /sbin/init
- Init - executes runlevel programs
- Runlevel - programs from /etc/rc.d/rcXX.d are started

2.4.5 BIOS <> UEFI

- Recent systems use UEFI, not BIOS
- UEFI is required to boot from GPT-disk
- Linux has no trouble working with UEFI

So why will we use MBR?

- Virtualisation is the norm
- Virtual machines typically have small disks
- Small disks are MBR partitioned

2.5 Filesystems

2.5.1 Windows

- FAT (1977)
- FAT32 (1996)
- NTFS (1993)
- ReFS (2012)

2.5.2 Linux

- Ext (1992)
- Ext2 (1993)
- Ext3 (2001)
- Ext4 (2008)
- ZFS (2005)
- Btrfs (2007)

2.5.3 Swap

= Paging

- Free up physical memory (RAM) by moving pages to slower storage (storage disks instead of RAM)
- Page out = memory page moves to swap
- "Swapiness"
 - = parameter between 0 and 100
 - = how quickly linux will swap
 - * 0 = very conservative
 - * 100 = very aggressive
- Windows uses a swap file (pagefile.sys)
- Linux uses a swap partition

2.6 File structure




Figure 3: Linux uses a tree structure




Figure 4: Windows uses a similar structure, but every volume uses a letter.




Figure 5: With linux, volumes are 'mounted' to folders somewhere under root /

2.7 Configuration

2.7.1 Packages

- Tools and applications are build up by files
- All files belonging to 1 application are bundled in a package
- Packages in debian have the .deb extension

Repositories

- Packages are collected in repositories
- Are made available through the internet
- Packages have dependencies

2.7.2 Package management

Debian: dpkg & apt (Advanced Package Tool)

- dpkg: Install, remove, give info about .deb packages
 - dpkg -l = lists packages
- apt: Get packages from a repository and install, remove, give info, ...
 - apt update
 - * Contact the repositories
 - * Get most recent list of packages and versions
 - apt upgrade
 - * Of the packages which are more recent in the repositories compared to what is installed: install newest version
 - apt install <xyz>
 - * Download package <xyz> from the repository

- * Check the dependencies and download depending packages
- * Install package <xyz> and all corresponding dependencies

Which repositories? See /etc/apt/sources.list for the list of repositories. You can add/remove/change repositories in this file.

2.7.3 Useful packages

- open-vm-tools
- vim
- sudo
- tcpdump

Install multiple packages in one command: `apt install vim sudo tcpdump ntp`

2.8 Shutdown of VM

- Power button (=ACPI shutdown)
- Shut down operating system only
 - = halt
- Shut down operating system and VM, multiple ways:
 - `shutdown -P now`
 - `init 0`
 - `poweroff`
- Reboot
 - `reboot`
 - `init 6`
 - `shutdown -r now`

2.9 Basic network

- No GUI ⇒
- Layer 1: Physical (VMWare virtual network)
- Layer 2: Datalink (Ethernet & MAC address)
- Layer 3: Network (IPv4)
- Layer 4: Transport (Transport Control Protocol (TCP), User Datagram Protocol (UDP))
- Layer 5: Application (SSH, HTTP, ...)

2.9.1 Basic networking commands

- arp
- ping
- route
- bmon

2.10 Services

- Processes that 'listen' on the network
 - TCP or UDP port
- Overview of currently running / listening services: ss command
 - ss -tulpn
 - t: show TCP
 - u: show UDP
 - l: show listening
 - p: show process ID
 - n: no name-resolving

2.11 Wooclap Questions

- Why do we talk about GNU/Linux?
- What is a kernel?
- What is the difference between Open Source and free?
- How is the Administrator user called? What is its uid?
- What is MBR?
- What are the limitations of MBR? (Solution?)
- What is swap? What is swappiness?
- What is a package?
- What is a repository?
- What is a dependency?
- What is a package manager?
- What is the difference between 'apt update' and 'apt upgrade'?
- Which protocol makes the link between MAC address & IP address?
- Which command gives you the current ARP-table?
- What are the 5 layers of the TCP/IP network model?
- How do you find the MAC-address of a network interface?
- Put Linux boot process in correct order (6 levels)

- What is a linux distribution?

3 File structure

- Tree structure
 - Leaves = files
 - Branches = directories
 - The tree is inverted, root = /
- Everything is a file (even devices, random numbers, and RAM) under 1 root
- This is in contrast to Windows, where every volume is a root.




Figure 6

3.1 Intermezzo: single user mode

- Linux (the kernel) is built up as a multi-user system from the beginning
- Standard behaviour = multi-user
- But: also possible to boot in single-user mode
 - No daemons, no multiple logins
 - Sometimes called **Maintenance mode**
- Examples of usage
 - Filesystem repairs
 - Upgrade of distribution
 - Password recovery
 - Adjustments to the root filesystem
 - Forensics after security incident

3.1.1 Runlevels

= predefined operating system status

- Is presented with a number
- Linux has 7 runlevels:
 - 0 = system halt (= VM shutdown)
 - 1 = single user
 - 2 = multi-user, no NFS (no network services, not often used)
 - 3 = multi-user, CLI (Command Line Interface)
 - 4 = self-definable
 - 5 = multi-user, GUI (Graphical User Interface, if installed)
 - 6 = reboot

3.2 Intermezzo: Add disk

Add a new disk without shutting down the system

1. Adjust VM: add disk
2. Detect added disk
3. Partition disk
 - fdisk (for MBR)
 - parted (for GPT)
4. Create filesystem
 - Partition = collection of blocks (sectors)
 - Not usable for the OS \Rightarrow create filesystem
 - `mkfs.ext4 /dev/sdb1`




Figure 7

5. Mount filesystem
 - `mkdir /mnt/datadisk`
 - `mount /dev/sdb1 /mnt/datadisk`
 - see if it worked: `df -h`

For detailed steps: see labs!

3.2.1 What after a reboot?

Use `/etc/fstab` = a file that contains what needs to be mounted at boot

- Device (`/dev/sdXY` or UUID)
- Mountpoint (`/mnt/folder`)
- Type of filesystem (`ext4`, `ntfs`, ...)
- Options

3.3 Navigate through the tree

- `pwd`
 - Print working directory
 - Shows where in the tree you are
- `ls`
 - Show a list of files in the working directory
 - `ls -la` : 10 characters at the beginning of each line. The `d ==` directory (see later)
- When you login, you are in your home directory
- `/` (= the filesystem root) is not the same as `/root` (the home directory of the root user)
- `.` = current directory
- `..` = the directory one higher

3.3.1 Relative vs absolute path

Relative paths:

- `cd ..` = go to the directory above the current directory
- `cat ../etc/issue` = go to the `etc` directory, one directory above the current directory. Open the `issue` file

Absolute paths:

- `cd /` = go to the root directory
- `cat /etc/issue` = go to the `etc/` directory under `/` (root)

3.4 Filesystem Hierarchy Standard (FHS)

- Describes how the filesystem in Linux is build up
- Maintained by the Linux Foundation
- Most recent version: v3.0 (2015)

3.4.1 Rules in the standard

- / is the root of the tree structure
- /bin
 - essential binaries (executable files), required for single user mode
- /boot
 - the place on the filesystem where the boot files reside
 - configuration files for GRUB
 - kernels
 - initrd
 - * initial RamDisk
 - * During boot a temporary root-filesystem is being created in RAM
 - * This is used so the kernel can load important modules, so it can then switch to the real root filesystem
 - * part of step 3 of the linux boot process (BIOS - MBR - GRUB - kernel - init - runlevel)
- /dev
 - Devices get a place in the filesystem
 - * sda
 - * rtc
 - * random
 - * cpu
 - * urandom
 - * null
 - ls -lah /dev/
- /etc
 - Host-specific system-wide configuration files
 - Configuration for this host, readable for the whole system
- /home
 - Each (non-system) user has a home directory
 - except for root \Rightarrow /root
- /mnt
 - (temporarily) 'mounted' filesystems
 - * Network shares
 - * USB-disk
 - * DVD-ROM
 - * Extra disks

- Some distributions use /media for this




Figure 8

- /opt
 - Optional application software packages
 - Our installation \Rightarrow no applications installed yet = empty (for now)
- /proc
 - Virtual filesystem
 - Provides information about processes and the linux kernel
 - cat /proc/cpuinfo
 - cat /proc/sys/net/ipv4/ip_forward
 - cat /proc/partitions
- /sbin
 - Essential system binaries
 - Only executable by root user
 - fsck, init, route
- /tmp
 - Directory for temporary files
 - Emptied at reboot (with most distributions)
- /usr
 - Read-only user data
 - Constains most user (non-root) utilities and applications
- /var
 - Variable files
 - Files that are expected to change continuously during normal system use
 - Logs, spool files, temporary e-mail files, ...

3.5 Some useful tips

3.5.1 History

```
1 ~# history
2
3 # shows a list of former commands executed by this user
4 # spans log-in sessions
5 # in reality, it shows the contents of the ~/.bash_history file
6 # if you use another shell like zsh, it's the ~/.zsh_history file
```

CTRL + r:

- Search the command history
- Show commands that match what you're typing
- repeatedly press ctrl+r to scroll through results

3.5.2 Bind mount

Situation

- /mnt/storage is the normal mountpoint for other filesystems (e.g. SAN)
- Filesystem could not be mounted, but a process already started writing data
- ⇒ this data arrives on the / filesystem under the directory /mnt/storage
- Problem fixed and filesystem can be mounted again ⇒ mounted under /mnt/storage
- ⇒ the already written data is now hidden

The solution

- Create /mnt/storage and put some data in it
- Create a 1GB disk, ext4 formatted, mount under /mnt/storage ⇒ data is now hidden
- Use mount -o bind to get data back without unmounting

3.5.3 dd

= Command to read or write bytes

```
1 # Example: overwrite first 2048 bytes of a disk with zeros
2 ~# dd if=/dev/zero of=/dev/sdb count=4 bs=512
3 # Example: overwrite disk with random data when taking out of service
4 ~# dd if=/dev/random of=/dev/sdb bs=1M
```

3.6 Wooclap Questions

- How do you ask the shell in which folder you are currently in?
- What is meant with the term 'runlevel' in Linux
- Describe single user mode with 1 word when you think of its primary use
- What is / are the most common runlevel(s) under linux? (So not all of them!)

- Where can you find the devices under Linux?
- What is the home directory of the root user?
- What command do we use to create a filesystem in a partition?
- What file do you need to edit to have a mounted filesystem available even after reboot?
- Where can you put temporary files in a linux system?
- How can you quickly search through your previously used commands?
- How do you quickly search through previously typed commands?
- What is swap?
- What are the limitations of MBR?
- What can you use a bind mount for?

4 Filesystems

4.1 Introduction

Books:

- A group of letters together = a word
- A group of words together = a sentence
- A group of sentences together = a book
- A collection of books together = a library
- Books are ordered/sorted according to a certain system
 - Best known: Dewey Decimal System

Computers:

- Work with 0's and 1's
- 1 character in ASCII or ISO-8859-1 = 8bits (1 byte)
- 1 Unicode character in UTF-8: between 8 and 32 bits (4 bytes)
- Gets stored on block devices
 - Hard devices, SSDs, RAMdisk, USB-stick
 - The opposite of block devices = character devices
- System needs to organize this

4.2 Blocks

- Disk = blocks
- Collection of blocks = sector (mostly 512 bytes)
- Collection of sectors = partition
- Partition not usable for an OS \Rightarrow filesystem needed

Definitie 4.1 (Filesystem) *A filesystem is the methods and data structures that an operating system uses to keep track of files on a disk or partition; that is, the way the files are organized on the disk.*

Several choices:

- Ext2/3/4
- BrtFS
- ZFS
- ...

4.3 ext2/3/4

Ext3 and Ext4 have journaling:

4.3.1 Journaling

- Keeping track of changes that have not been committed to disk in a sort of 'diary'
- A kind of logbook of previous actions
- Why?
 - Bring filesystem online faster after system crash or power failure

4.4 RAID

Definitie 4.2 (RAID) *Redundant Array of Independant Disks (RAID) is a data storage virtualisation technology that combines multiple physical disk drive into one ore more logical units.*

Many purposes:

- *Data redundancy*
- *Performance Improvement*
- *Both*

4.4.1 RAID Controller

- Disks are connected to the controller
- The RAID controller displays the disks as 1 disk to the OS
- Nowadays, we call the RAID Controller the Host Bus Adapter (HBA)
-

4.4.2 RAID 0

RAID level 0 uses striping:

Definitie 4.3 (Striping) *Data striping is the technique of segmenting logically sequential data (files) so that segments are stored on different physical storage devices*

Purpose:

- *Increasing data throughput*

- *Balancing I/O load accross an array of disks*




Figure 9: Example: files A and B (4 blocks each) are spread over disks D1-D3

4.4.3 RAID 1

RAID level 1 uses mirroring:

Definitie 4.4 (Mirroring) *Disk mirroring is the replication of logical disk volumes onto seperate physical disks.*

Purpose:

- *Continuous availability: in case of hardware failure, you always have a backup of your data*
- *Increasing read speeds*




Figure 10: RAID 1

4.4.4 RAID 4

- If we have at least 3 disks
- For every block of data:
 - Divide the block in 2 halves: A and B
 - Write A to disk 1
 - Write B to disk 2
 - Write A+B to disk 3
- \Rightarrow RAID 4 is striping (disk 1 & 2) with parity (disk 3)
- Capacity x2
- Read speed x2
- Write speed is limited, because of the need to write all parity data to a single disk




Figure 11: A RAID-4 setup with 4 disks. Disk 3 is the parity disk

4.4.5 RAID 5

RAID level 5 like RAID 4, but the parity is distributed.

- This evens out the stress of a dedicated parity disk (RAID 4)
- Write performance is increased




Figure 12: RAID-5: distributed parity with 4 disks

4.4.6 RAID 6

RAID level 6 like RAID 5, but with a second parity block.




Figure 13

4.4.7 Disk failure

For every RAID level, a certain amount of disks can fail without it becoming a problem:

- RAID 0: no disks can fail: if any disk fails, you lose data
- RAID 1: Every disk except for one can fail
- RAID 5: 1 disk can fail
- RAID 6: 2 disks can fail

4.4.8 Compound RAID levels

Combining RAID levels is possible:

- RAID 10 = RAID 1 + RAID 0
- RAID 01 = RAID 0 + RAID 1
- RAID 50 = RAID 5 + RAID 0

4.5 OpenZFS

4.5.1 ZFS

- Zettabyte File System
- Developed by Sun (2011) \Rightarrow Open source
- Now Oracle (2010) \Rightarrow Not free and closed source

4.5.2 Open ZFS

- Fork van ZFS
- 2013
- Option in Ubuntu-installer

Features:

- Long term storage
- Checksum of all data and metadata
- Native RAID levels (0, 1, 5, 6, ...)
- All data gets written through Copy-On-Write:
 - COW = when a write request is made, the data is copied into a new storage area, and then the original data is modified.
 - Redirect-on-write or ROW: the original storage is never modified. When a write request is made, it is redirected away from the original data into a new storage area.
- Snapshots (read-only and mountable)
- Transparent compression
- Huge storage possibilities: up to 256 quadrillion zettabytes
- 128 bits system

4.6 Intermezzo: Kernel modules

- Linux = kernel
- Kernel = modular
- /boot/config-4.9.0-13-amd64: config for this kernel
 - Describes what is inside this kernel
- Not all modules are loaded all the time

4.6.1 Commands

```
1 # Request current list of modules:
2 ~# lsmod
3
4 # Load module:
5 ~# modprobe brtfs
6
7 # Remove module ("unload"):
8 ~# rmmod brtfs
```

4.7 Intermezzo: Snapshots

- Literally: a photograph of your filesystem
- Captures the state of the filesystem at a certain point in time
- "The possibility to return in time"

4.7.1 Do we still need backups if we have snapshots?

YES!

- RAID 1 (mirroring) only protects against disk failure, nothing else
- If someone deletes all data from one disk, the RAID controller will delete all data from the other disk.
- Snapshots can get lost: what if your server fails?
- ⇒ backups can be stored safely, on other disks

5 File manipulation

5.1 Basics

```
1 # create an empty file called 'test'
2 ~$ touch test
3
4 # edit a file
5 ~$ vim test
6
7 # remove file
8 ~$ rm rabbit
9
10 # move the file to /tmp
11 ~$ mv test /tmp/
12
13 # rename the file
14 ~$ mv test rabbit
15
16 # Linux doesn't really look at file extensions
17 # check the file extension:
```

```

18 ~$ file <filename>
19 ~$ file /boot/inird.img-4.9.0-13-amdb64
20 ~$ file /etc/init.d/networking

```

5.2 Bundle files

- Tape ARchiver: TAR
 - Created originally to bundle files/directories for storage on tapes
- You can combine tar with gzip: .tar.gz
 - tar cfv bundle.tar *.txt ⇒ not compressed
 - tar czfv bundle.tar.gz ⇒ compressed

```

1 ~$ mkdir bundle
2 ~$ cd bundle
3 ~$ touch 1.txt 2.txt 3.txt
4 ~$ tar cfv bundle.tar *.txt
5 # c = create a new archive
6 # f = specify a filename (bundle.tar)
7 # v = verbose: show what happens
8 ~$ tar --list -f bundle.tar
9 ~$ file bundle.tar
10
11 # extracting
12 ~$ tar zxvf bundle.tar.gz
13 # z = zipped (compressed)
14 # x = eXtract
15 # v = verbose
16 # f = the argument (a file)

```

5.3 Links and inodes

- Modern filesystems support links
- This is different from shortcuts in Windows:
- Windows shortcuts are text files that refer to other files

5.3.1 Inodes

Definitie 5.1 (Inode) *An inode is a data structure on a filesystem on Unix-like operating systems that stores all the information about a file except its name and its actual data*

Metadata: data about the file

- Creation date
- Creation author
- Access rights
- ...

5.3.2 Symbolic links

Definitie 5.2 A symbolic link (also symlink or soft link) is a term for any file that contains a reference to another file or directory in the form of an absolute or relative path

Also called 'softlinks'

```
1 ~$ ln -s <target> <link-name>
2 ~$ ln -s /etc/issue test-link
3 # try out the following commands after creating a link:
4 ~$ cat test-link
5 ~$ file test-link
6 ~$ cat /etc/issue
```

5.3.3 Hardlinks

- Same file, different name
- A hardlink refers to an inode, while a softlink refers to a file (which refers to an inode)




Figure 14: Symlink vs Hardlink

5.4 File permissions




Figure 15: The output of 'ls -l' creates this type of output

1. first character: type of file (directory, symbolic link, block device, character device)

2. next 9 characters: owner rights, group rights, other rights
3. next number: number of references to the inode
4. owner user
5. owner group
6. size of file
7. last modified
8. name of file

```

1  # change the owner and group of a file or directory
2  ~# chown <user>:<group> <file>
3  ~# chown root:staff file.txt
4
5  # change the rights for a file
6  # chmod: change mode

```

```

0 --- indicates no permissions
1 --x indicates execute permissions
2 -w- indicates write permissions
3 -wx indicates write and execute permissions
4 r-- indicates read permissions
5 r-x indicates read and execute permissions
6 rw- indicates read and write permissions
7 rwx indicates read, write, and execute permissions

```

Figure 16: Octal notation

5.5 Overview of basic commands

Usage and details for these commands: see labs

- cat: print contents of file to terminal
- cut: cut (structured) input on a specific place: show a certain column, etc. . .
- grep: display lines for which the pattern matches
- egrep: extended grep, better handling of regular expressions
- find: search for files in a hierarchy of files and directories
- head: show first lines of file
- tail: show last lines of file
- less: show the contents of a text file, interactively
- man: show manual page for specific command
- wc: word count (but also character count, byte counts, newline counts, . . .) for a file
- date: show or configure system date and time
- cal: show a textual calendar
- sort: sort a file
- uniq: in a sorted output: count double lines or only show unique lines

5.6 Wooclap

- What is meant by the term journaling for filesystems?
- Why is journaling used with filesystems?
- Give 2 examples of filesystems under linux that use journaling.
- How can you find out which kernel modules are currently loaded?
- Which command can you use to load a kernel module?
- And which to 'unload' a kernel module?
- How many disks do you need at least to build a RAID10 system? Why?
- What is meant by a 'Copy On Write' filesystem?
- What are the advantages of a CoW filesystem?
- What are snapshots (in the context of storage systems)?
- What are the disadvantages of a CoW filesystem?
- Why do you still need backup when you have RAID1 and have snapshots?
- How can you find out which 'type' is a file? There are no extensions.
- What is an inode?
- What is the difference between a softlink and a hard link?
- At the output of the command `ls -la`: Which values can the first character of the line have and what do they mean?
- At the output of the command `ls -la`: Which possible values can the 3 groups of 3 characters have to describe the rights?
- With what command can you 'change' the 'owner' of a file or directory?
- With which command can you 'change' the rights of a file or directory?
- What does number 5 mean when you use it to determine file system permissions?
- What does number 7 mean when you use that to determine file system rights? Explain why.
- Which command do you use to cut structured input at a specific location?
- Which command do you use to display the first 16 lines of a text file?
- Which command can you use to find out all the modified files from the last 24 hours?
- Which command do you use to display the last 12 lines of a text file?
- How can you find out how long it has been since a linux system was rebooted?
- Which command can you use to get an overview of all daemons that are currently active in your system?

6 Text editors, Piping, Redirection & Jobs

6.1 Text editors

- Emacs (productivity, extensibility)

- Nano (simplicity)
- Vi / Vim (=VI iMproved) (productivity)
- Ne

Our choice: Vim

6.1.1 vi vs vi-improved

- Navigating in vi: HJKL (left, down, up, right)
- Navigating in vim: HJKL or arrow keys




Figure 17: Everything you can do in vi, you can do in vim, and more!

6.1.2 First steps in vim

```

1  # start vim:
2  $~ vim
3
4  # start vim tutorial:
5  $~ vimtutor

```

- bottom left of window: the current mode
- INSERT = the mode that lets you enter text
- Enter INSERT mode: i
- Leave INSERT mode and go to normal mode: ESC
- Once in normal mode, you can enter commands using ':'
- ESC : q ⇒ quit
- ESC : w ⇒ quit
- ESC : wq ⇒ write and quit

```

1  # copy a line of text
2  ESC
3  # put cursor on the line you want to copy
4  yy # yank yank: copy a line of text

```



```

5 | p # put: paste the copied line
6 |
7 | # copy 2 lines of text and paste it 8 times:
8 | ESC 2yy # 2 yank yank: yank 2 lines
9 | 8p      # 8 put: paste the copied lines 8 times

```

6.1.3 Search and replace

You can easily search and replace in a text file, even with regex:

```

1 | # search and replace the next instance:
2 | :s/old/new
3 |
4 | # all instances: TODO: difference between :s and :%s
5 | :s/old/new/g
6 |
7 | # all instances between line x and y:
8 | :x,ys/old/new/g
9 |
10 | # all instances in a complete file:
11 | :%s/old/new/g
12 |
13 | # all instances in whole file, with confirmation:
14 | :%s/old/new/gc
15 |
16 | #undo:
17 | (ESC) u

```

6.1.4 Basic editing tricks

```

1 | # starting from line 4, indent the next 7 lines:
2 | (ESC) 7>>
3 |
4 | # remove indentation on line 10
5 | (ESC) 10gg
6 | <<
7 |
8 | # delete line 3
9 | (ESC) 3gg
10 | dd
11 |
12 | # delete the next 4 lines
13 | (ESC) 4dd
14 |
15 | # enable and disable syntax highlighting in vim
16 | (ESC) : syntax on
17 | (ESC) : syntax off

```

For more tricks: see labs

6.2 Piping

= Use the output from one command as input for the next command

```
1 # sort the music file alphabetically and count the number of occurrences of each unique line
2 sort music.txt | uniq -c
3
4 # count the number of unique lines in music.txt
5 sort music.txt | uniq | wc -l
6
7 # count the number of lines in /etc/locale.gen where nl or NL occurs
8 grep -i nl /etc/locale.gen | wc -l
9
10 # count the number of lines in /var/log/syslog where kernel occurs
11 cat /var/log/syslog | grep kernel | wc -l
12
13 # count the number of lines in /var/log/syslog where kernel does NOT occur
14 cat /var/log/syslog | grep -v kernel | wc -l
15
16 # Show of what days there are logs in /var/log/syslog
17 # the sixth field is the day field:
18 cat /var/log/syslog | cut -c 6 | uniq
19
20 # show the different sources of log entries in /var/log/syslog
21 # kernel, client, systemd, ...
22 cat /var/log/syslog | cut -d' ' -f5 | cut -d '[' -f1 | sort | uniq -c
```

6.3 Redirection

Do not send the output of a command to stdout, but to another location, like a textfile

```
1 # to overwrite a file (or create if it doesn't exist)
2 ls -la > listing.txt
3
4 # to append to a file (or create if it doesn't exist)
5 ls -la >> listing.txt
6
7 # these two commands have the same result
8 cat > textfile.txt
9 touch textfile.txt
10
11 # redirect the output of 'ls -la' to a file with a custom name:
12 # example: output_2021-03-03
13 ls -la > output_$(date +%F)
```

6.3.1 stdout and stderr

= 2 important output streams

- Normal situation: stdout and stderr appear on the terminal
- Redirection: stdout to a file

- stderr still prints to the terminal
- Redirect stderr: 2> errorfile.txt



Figure 18

```

1 # redirect the output of a command to out.txt
2 # and redirect the error of the command to error.txt:
3 ls -la > out.txt 2> error.txt
4
5 # redirect stderr to stdout (01), and then redirect stdout to a file out.txt:
6 ls -la > out.txt 2>&1
7
8 # redirect both to a file:
9 ls -la &> out

```

6.3.2 stdin

```

1 # this command prints the amount of lines in a file
2 wc -l music.txt
3
4 # this command does the same
5 wc -l < music.txt
6
7 # this command does the same, but prints the output of wc to out.txt
8 wc -l < music.txt > out.txt

```

6.4 Jobs and process Management

When you execute a command: a process is started

- Every process gets a process ID (PID)
- The init process has PID 0. It starts other processes.
- Every process has a parent
 - ⇒ tree structure of processes
 - Get insights into this structure with 'pstree' (part of the 'psmisc' package)
 - Process stops: exit code is passed to the parent

```

1 # report a snapshot of current processes
2 ps
3

```

```

4 # display a tree of processes
5 pstree -p

```

6.4.1 Exit codes

```

1 # when a process stops:
2 # if bash was parent => exit code is passed to bash
3 # exit code is available in the $? variable:
4 # read contents of a variable with echo:
5 echo $?
6 0
7
8 # 0 = ended successfully without errors
9
10 # 1 = not ended successfully, there were errors

```

```

1 # test = verify if you have rights on a file (with -r = read rights)
2 ~$ test -r music.txt
3
4 ~$ echo $?
5 0
6
7 ~$ test -r doesnotexist
8
9 # the file doesn't exist, so it exited with code 1:
10 ~$ echo $?
11 1

```

6.4.2 Combining commands

= not the same as piping!

```

1 # note the difference between:
2 # & = execute command 1, then execute command 2
3 # && = execute command 1, and only execute command 2 if command 1 was successful (exit code = 0)
4
5 # one ampersand
6 test -r test.txt & echo "MCT rocks"
7
8 # two ampersands
9 ~# test -r test.txt && echo "MCT rocks"
10
11 ~# test -r doesnotexist & echo "MCT rocks"
12
13 # the first command will exit with code 1
14 # so the second command will not execute
15 ~# test -r doesnotexist && echo "MCT rocks"

```

6.4.3 Jobs

A job is a new process originating from the same parent

```
1 # start a command as job
2 tail -f /var/log/syslog &
3 # output appears on stdout, but process runs as a job in the background
4
5 # bring a job to the foreground
6 fg <index>
7
8 # stop the job, but do not terminate:
9 CTRL-Z
```

6.4.4 Inter-process Communication

A signal is an asynchronous notification sent to a process or thread within that process to notify that there has been an event

Signal sent to process \Rightarrow OS interrupts normal execution of that process to deliver the signal

Sending a signal to a process: with the 'kill' command:

```
1 # not only to kill a process, also to send other signals
2 kill -s <signal>
```

Signals

- SIGHUP - 1 - terminate (hang up)
- SIGINT - 2 - terminal interrupt signal
- SIGKILL - 9 - kill (cannot be caught or ignored)
- SIGTERM - 15 - termination signal

```
1 # send signal 15 to a process with the entered PID
2 kill -s 15 <pid>
3
4 # send signal 15 to the PID of the tail process
5 kill -s 15 'pidof tail'
6
7 # send signal 9 to a process with the entered PID
8 kill -s 9 <pid>
9
10 # kill the process with name 'tail'
11 pkill tail
```

6.5 Intermezzo: System Load

= a number which represents the load on a computer system

- Completely idle system: system load 0
- Each process which uses a resource or is waiting for a resource: system load + 1

- Gives an indication of how heavy a computer system is loaded
- System load is a snapshot, doesn't say anything
 - System load of 17: is that a problem? No.
 - More interesting: the evolution of the systemload over time

```

1  # show the system load of the last minute, last 5 minutes and last 15 minutes:
2  uptime
3
4  # show who is logged on and what they are doing
5  w
6
7  # display linux processes
8  top
9  # or better:
10 htop

```

6.6 Some useful tips

6.6.1 With which unique IP-addresses are there open sockets and how many?

```

1  netstat -anpt | awk '{print $5}' | sort | uniq -c

```

6.6.2 TTY

= Tele TYpewriter = a terminal which is connected with stdin

```

1  # print the filename of the terminal currently connected to standard input:
2  ~$ tty

```